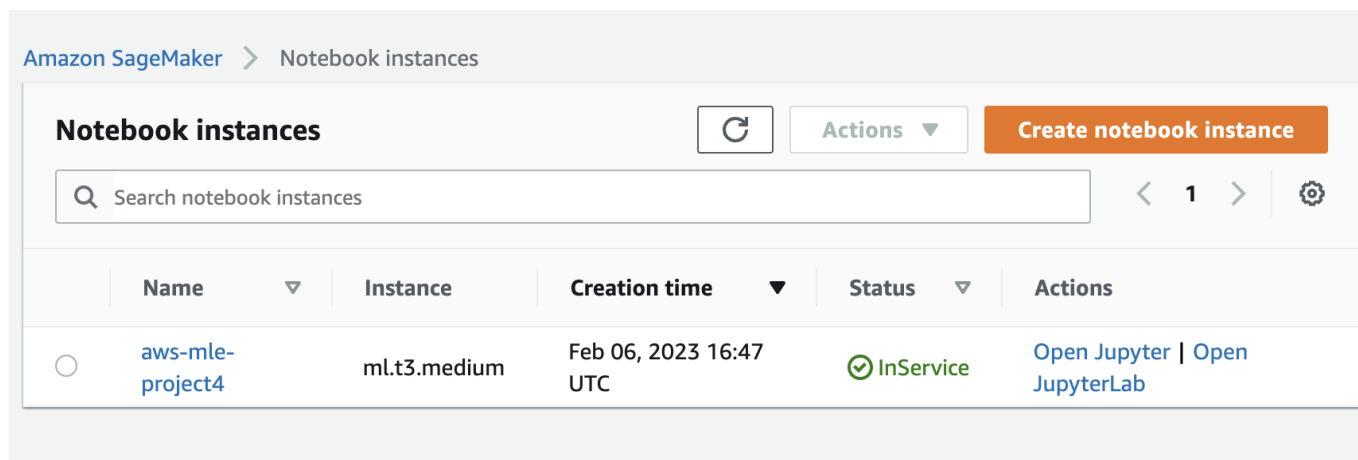







# Operationalizing an AWS Machine Learning Project





## Initial Set Up



Amazon SageMaker > Notebook instances

**Notebook instances**  **Actions**  **Create notebook instance**

 **1**  

	Name 	Instance	Creation time 	Status 	Actions
<input type="radio"/>	aws-ml-project4	ml.t3.medium	Feb 06, 2023 16:47 UTC	 InService	<a href="#">Open Jupyter</a>   <a href="#">Open JupyterLab</a>

The ml.t3.medium is the cheapest among Sagemaker's standard instances at \$0.05 per hour. It has 5 GiB of memory and runs on 2 vCPU. Optimized to start within 2 minutes, this instance would mean less waiting time in between starting and stopping the instance. This is important as one way of sticking to a limited budget is to work remotely and locally simultaneously, requiring starting and ending the instance repetitively.

Although there are other instance types with the same fast launch feature, like the more powerful ml.g4dn.xlarge, which allows GPU-based capabilities, it is best to start conservatively with a smaller instance before moving on to a bigger instance should a faster performance be required for the workload, considering that this is a small, personal project.

## Initial Training & Deployment

```
hyperparameters = {'batch_size': 64, 'learning_rate':  
'0.037260043722494224'}  
estimator = PyTorch(  
    entry_point='hpo.py',  
    base_job_name='dog-pytorch',  
    role=role,  
    instance_count=1,  
    instance_type='ml.m5.xlarge',  
    framework_version='1.4.0',  
    py_version='py3',  
    hyperparameters=hyperparameters,  
    ## Debugger and Profiler parameters  
    rules = rules,  
    debugger_hook_config=hook_config,  
    profiler_config=profiler_config,  
)
```

```
INFO:botocore.credentials:Found credentials from IAM Role: BaseNotebookInstanceEc2InstanceRole
INFO:sagemaker:Creating model with name: pytorch-inference-2023-02-12-15-28-21-386
INFO:sagemaker:Creating endpoint-config with name pytorch-inference-2023-02-12-15-28-21-903
INFO:sagemaker:Creating endpoint with name pytorch-inference-2023-02-12-15-28-21-903
-----!
```

	Name ▾	ARN	Creation time ▾	Status ▾	Last updated
○	<a href="#">pytorch-inference-2023-02-12-15-28-21-903</a>	arn:aws:sagemaker:us-east-1:663876033295:endpoint/pytorch-inference-2023-02-12-15-28-21-903	Feb 12, 2023 15:28 UTC	✔ InService	Feb 12, 2023 15:30 UTC

Endpoint settings

Name

pytorch-inference-2023-02-12-15-28-21-903

ARN

arn:aws:sagemaker:us-east-1:663876033295:endpoint/pytorch-inference-2023-02-12-15-28-21-903

Status

✔ InService

Creation time

Sun Feb 12 2023 16:28:22 GMT+0100 (Central European Standard Time)

Type

Real-time

Last updated

Sun Feb 12 2023 16:30:52 GMT+0100 (Central European Standard Time)

URL

https://runtime.sagemaker.us-east-1.amazonaws.com/endpoints/pytorch-inference-2023-02-12-15-28-21-903/invocations

[Learn more about the API](#)

## Multiple Instance Training & Deployment

The time it took for the model to be trained was no different from that of the singular instance, however, the model with 3 instances did a better job at classifying the image, considering that the actual image has a label of 11.

```
mi_estimator = PyTorch(
    entry_point='hpo.py',
    base_job_name='multi-dog-pytorch',
    role=role,
    instance_count=3,
    instance_type='ml.m5.xlarge',
    framework_version='1.4.0',
    py_version='py3',
    hyperparameters=hyperparameters,
    ## Debugger and Profiler parameters
    rules = rules,
```

```
        debugger_hook_config=hook_config,
        profiler_config=profiler_config,
    )
```

	Name ▾	Creation time ▼	Duration	Job status ▾	Warm pool status	Time left
<input type="radio"/>	<a href="#">multi-dog-pytorch-2023-02-12-16-10-23-699</a>	Feb 12, 2023 16:10 UTC	21 minutes	✔ Completed	-	-
<input type="radio"/>	<a href="#">dog-pytorch-2023-02-12-14-45-50-440</a>	Feb 12, 2023 14:45 UTC	21 minutes	✔ Completed	-	-

```
INFO:botocore.credentials:Found credentials from IAM Role: BaseNotebookInstanceEc2InstanceRole
INFO:sagemaker:Creating model with name: pytorch-inference-2023-02-12-16-41-31-338
INFO:sagemaker:Creating endpoint-config with name pytorch-inference-2023-02-12-16-41-31-839
INFO:sagemaker:Creating endpoint with name pytorch-inference-2023-02-12-16-41-31-839
-----!
```

### Endpoint settings

Name	Type
pytorch-inference-2023-02-12-16-41-31-839	Real-time
ARN	Last updated
arn:aws:sagemaker:us-east-1:663876033295:endpoint/pytorch-inference-2023-02-12-16-41-31-839	Sun Feb 12 2023 17:43:46 GMT+0100 (Central European Standard Time)
Status	URL
✔ InService	https://runtime.sagemaker.us-east-1.amazonaws.com/endpoints/pytorch-inference-2023-02-12-16-41-31-839/invocations
Creation time	<a href="#">Learn more about the API</a>
Sun Feb 12 2023 17:41:32 GMT+0100 (Central European Standard Time)	

```
pred[0]
# single instance => 28
# multiple instance => 11
```

## EC2 Training

Unlike the demo shown in the module, there were only 3 choices for Deep Learning AMIs and only 2 of them have an environment that supports Pytorch packages and dependencies.



```

hpo.py 9+ X
hpo.py > ...
145     logger.info("Saving Model")
146     torch.save(model.cpu().state_dict(), os.path.join(args.model_dir, "model.pth"))
147
148     if __name__ == '__main__':
149         parser=argparse.ArgumentParser()
150         parser.add_argument('--learning_rate', type=float)
151         parser.add_argument('--batch_size', type=int)
152         parser.add_argument('--data', type=str, default=os.environ['SM_CHANNEL_TRAINING'])
153         parser.add_argument('--model_dir', type=str, default=os.environ['SM_MODEL_DIR'])
154         parser.add_argument('--output_dir', type=str, default=os.environ['SM_OUTPUT_DATA_DIR'])
155
156         args=parser.parse_args()
157         print(args)
158
159         main(args)
160
ec2train1.py 9+ X
ec2train1.py > ...
132     return train_data_loader, test_data_loader, validation_data_loader
133
134     batch_size=2
135     learning_rate=1e-4
136     train_loader, test_loader, validation_loader=create_data_loaders('dogImages',batch_size)
137     model=net()
138
139     criterion = nn.CrossEntropyLoss()
140     optimizer = optim.Adam(model.fc.parameters(), lr=learning_rate)
141
142     logger.info("Starting Model Training")
143     model=train(model, train_loader, validation_loader, criterion, optimizer)
144     torch.save(model.state_dict(), 'TrainedModels/model.pth')
145     print('saved')
146
147

```

The main difference between the hpo.py script and the ec2train1.py script lies in the way the arguments and hyperparameters are introduced to the model. In the ec2train1.py script, the hyperparameters and arguments are declared at the end of the script whereas the arguments and hyperparameters of the hpo.py script are declared within the Sagemaker notebook instance and are introduced to the script using `parser=argparse.ArgumentParser()` and a main function.

```

if __name__ == '__main__':
    parser=argparse.ArgumentParser()
    parser.add_argument('--learning_rate', type=float)
    parser.add_argument('--batch_size', type=int)
    parser.add_argument('--data', type=str,
default=os.environ['SM_CHANNEL_TRAINING'])
    parser.add_argument('--model_dir', type=str,
default=os.environ['SM_MODEL_DIR'])

```

```
parser.add_argument('--output_dir', type=str,
                    default=os.environ['SM_OUTPUT_DATA_DIR'])

args=parser.parse_args()
print(args)

main(args)
```

## Lambda

Lambda, Amazon's serverless compute service, is the ideal solution for small tasks that are frequently used as it executes code without underlying infrastructures like operating system or hardware specifications that can sometimes impede the smooth implementation of programs.

It is developed using Python code through Boto3, an AWS SDK that allows the function to interact with and manage AWS services, provided that it has the correct policies and execution role. This is done through the client as specified in the code shown below:

```
runtime=boto3.Session().client('sagemaker-runtime')
```

The code is executed by the handler function which is usually contained in a file called `lambda_function.py`. When a payload in the form of a JSON object is delivered, the Lambda function executes the code defined by the endpoint using the `invoke_endpoint()` method.

## Security and Testing a Lambda Function

Once the Lambda function has been written and deployed, a test case can be configured by creating a JSON object that matches which contains the arguments specified in the function. As AWS Lambda is provided with its own execution role, it is imperative that the correct policies are attached to it so that the one can test the function successfully.

IAM > Roles > pytorch-dogs-role-ursdoukz

pytorch-dogs-role-ursdoukz

Delete

Summary

Edit

Creation date

February 14, 2023, 17:41 (UTC+01:00)

Last activity

None

ARN

arn:aws:iam::663876033295:role/service-role/pytorch-dogs-role-ursdoukz

Maximum session duration

1 hour

Permissions

Trust relationships

Tags

Access Advisor

Revoke sessions

Permissions policies (2) Info

You can attach up to 10 managed policies.

Filter policies by property or policy name and press enter.

1

Policy name

Type

Description

AWSLambdaBasicExecutionRole-a48f462f-cb79-4886-9e4d-05e6...

Customer managed

AmazonSageMakerFullAccess

AWS managed

Provides full access to Amazon Sage

If this is done correctly, it will result to a response with a status code of 200 and the values specified in the return statement will be found in the body of the response as shown below.

Execution result: +

Status: Succeeded

Max memory used: 75 MB

Time: 1343.56 ms

Test Event Name

test\_dog

Response

```
{
  "statusCode": 200,
  "headers": {
    "Content-Type": "text/plain",
    "Access-Control-Allow-Origin": "*"
  },
  "type-result": "<class 'str'>",
  "Content-Type-In": "<__main__.LambdaContext object at 0x7f3958fc3c40>",
  "body": "[[0.22258417308330536, 0.2867152690887451, 0.23660334944725037, 0.3973812758922577, 0.5946304798126221, 0.32388997077941895, 0.14295358955860138, 0.2593267
```

Function Logs

START RequestId: b7fdccd8-6ade-4b0f-adae-c110bdf083b8 Version: \$LATEST  
Context::: <\_\_main\_\_.LambdaContext object at 0x7f3958fc3c40>  
EventType:: <class 'dict'>  
END RequestId: b7fdccd8-6ade-4b0f-adae-c110bdf083b8  
REPORT RequestId: b7fdccd8-6ade-4b0f-adae-c110bdf083b8 Duration: 1343.56 ms Billed Duration: 1344 ms Memory Size: 128 MB Max Memory Used: 75 MB

Request ID

b7fdccd8-6ade-4b0f-adae-c110bdf083b8

```
Test Event Name
test_dog

Response
{
  "statusCode": 200,
  "headers": {
    "Content-Type": "text/plain",
    "Access-Control-Allow-Origin": "*"
  },
  "type-result": "<class 'str'>",
  "Content-Type-In": "<__main__.LambdaContext object at 0x7f3958fc3c40>",
  "body": "[[0.22258417308330536, 0.2867152690887451, 0.23660334944725037, 0.3973812758922577, 0.5946304798126221, 0.32388997077941895, 0.14295358955860138, 0.2593267
```

7 / 9

```
"type-result": "<class 'str'>",
"CContent-Type-In": "<__main__.LambdaContext object at 0x7f3958fc3c40>",
"body": "[[0.22258417308330536, 0.2867152690887451, 0.23660334944725037,
0.3973812758922577, 0.5946304798126221, 0.32388997077941895,
0.14295358955860138, 0.2593267858028412, -0.2645488977432251,
-0.0056010279804468155, 0.3683377206325531, 0.40309959650039673,
-0.011438594199717045, 0.3243323266506195, 0.4832281768321991,
0.11940720677375793, 0.3301726281642914, 0.016071753576397896,
0.13358867168426514, 0.4334683120250702, 0.3398759961128235,
-0.12303052097558975, 0.38373783230781555, 0.208574116230011,
-0.17253750562667847, -0.13229110836982727, 0.41559934616088867,
-0.3317200839519501, 0.5983749032020569, 0.15570591390132904,
0.2418746054172516, 0.5634035468101501, -0.06630165129899979,
0.260342001914978, 0.15422964096069336, 0.26249510049819946,
0.05999879539012909, 0.17178316414356232, 0.28845247626304626,
0.14137883484363556, 0.2950528860092163, 0.34624922275543213,
0.1801588386297226, 0.37983566522598267, 0.1358564794063568,
0.39831653237342834, 0.09975286573171616, 0.06437289714813232,
0.20123106241226196, 0.2751830816268921, 0.3577496409416199,
0.10020403563976288, 0.09287401288747787, 0.2717002034187317,
0.10770490020513535, 0.19640129804611206, 0.40510982275009155,
0.02587900683283806, -0.0023690317757427692, 0.05544339120388031,
0.2715991735458374, 0.0018166087102144957, 0.1619286835193634,
-0.12189028412103653, 0.04019118845462799, -0.28905490040779114,
-0.18545860052108765, 0.35718852281570435, 0.02673361450433731,
0.06363802403211594, 0.3299733102321625, 0.10386842489242554,
-0.2123139351606369, -0.02947426214814186, 0.026548288762569427,
0.33504170179367065, -0.09340201318264008, -0.23949125409126282,
0.19283726811408997, -0.06110447272658348, -0.06545177847146988,
0.1155146136879921, 0.05927373468875885, 0.28965985774993896,
-0.19908912479877472, 0.05027751624584198, 0.3238394558429718,
0.20499038696289062, 0.015548424795269966, 0.2560834586620331,
0.25757649540901184, 0.024451695382595062, -0.1503819227218628,
-0.037278901785612106, 0.10541312396526337, 0.116038016974926,
-0.05723908543586731, -0.04249229282140732, -0.11206389963626862,
-0.2917730212211609, -0.01494930312037468, -0.32105588912963867,
0.3950178921222687, -0.3360593318939209, -0.322693407535553,
0.0007147123105823994, -0.07922632992267609, -0.4352700412273407,
-0.19211743772029877, -0.1981441080570221, -0.08153677731752396,
0.18471316993236542, -0.0832882970571518, -0.27037137746810913,
0.3735558092594147, -0.366347998380661, 0.0038887872360646725,
0.2003762274980545, -0.21762414276599884, -0.048697978258132935,
-0.3799019455909729, -0.37747669219970703, -0.14877746999263763,
0.009409474208950996, -0.2609257102012634, -0.37330231070518494,
-0.05776982381939888, -0.342445969581604, 0.03831148520112038,
0.062469542026519775, -0.41384974122047424, -0.5424782037734985,
-0.3256552815437317]]]"
}
```

As this is a course activity that does not involve private information, it does not require as much security as a project within a company. In this case, assigning any policy with full access, despite the requirements that it may fulfill in the execution of the code, can cause security issues. In this case, the additional protection by



segmentation provided by the VPC can be useful. After all, giving the minimal access that is required to accomplish tasks in a certain role will be the best way to maintain a secure system.

## Concurrency and Autoscaling

Concurrency decreases the latency of response during high-traffic situations. As this is a course project, it is quite easy to foresee that the function will not be exposed to a high amount of traffic. Needless to say that the traffic will be under the control of the coder, so a low value of 4 reserved concurrencies was chosen. With the project's low budget and very predictable traffic, there was really no need for a provisioned concurrency.

The same considerations applied to the autoscaling configuration. It does not require a high target value as this was done for the purpose of a small-scale project. In this case, it was best to choose 3 instances as the maximum and to lower the cost, a target value of 50 was chosen. As for response time, 30 seconds was chosen for both scale in and scale out times.

This means that 30 seconds after the number of simultaneous invocations reaches 50, the instances will increase and will decrease after the same amount of time once number of invocations decrease. The shorter scale in and scale out time in this case increases the responsiveness which makes up for the higher target value.